

```
In [1]: import pandas as pd
import numpy as np
df = pd.read_csv('Students Performance .csv')

df.head(10)
```

Out[1]:

	Student_ID	Student_Age	Sex	High_School_Type	Scholarship	Additional_Work	Spo
0	STUDENT1	19-22	Male	Other	50%	Yes	
1	STUDENT2	19-22	Male	Other	50%	Yes	
2	STUDENT3	19-22	Male	State	50%	No	
3	STUDENT4	18	Female	Private	50%	Yes	
4	STUDENT5	19-22	Male	Private	50%	No	
5	STUDENT6	19-22	Male	State	50%	No	
6	STUDENT7	18	Male	State	75%	No	
7	STUDENT8	18	Female	State	50%	Yes	
8	STUDENT9	19-22	Female	Other	50%	No	
9	STUDENT10	19-22	Female	State	50%	No	

```
In [2]: print(df.columns.tolist())
```

['Student_ID', 'Student_Age', 'Sex', 'High_School_Type', 'Scholarship', 'Additional_Work', 'Sports_activity', 'Transportation', 'Weekly_Study_Hours', 'Attendance', 'Reading', 'Notes', 'Listening_in_Class', 'Project_work', 'Grade']

```
In [7]: # student_performance_prediction_updated.py
```

```
import pandas as pd
import numpy as np
import joblib
import os
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler, OrdinalEncoder
from sklearn.linear_model import LinearRegression, Ridge, Lasso, LogisticRegression
from sklearn.ensemble import RandomForestRegressor, RandomForestClassifier, GradientBoostingClassifier
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score, accuracy_score
import warnings
warnings.filterwarnings('ignore')

class StudentPerformancePredictor:
    def __init__(self):
        self.model = None
        self.scaler = StandardScaler()
        self.label_encoders = {}
        self.grade_encoder = None
        self.feature_columns = None
```

```

        self.is_trained = False
        self.problem_type = None # 'regression' or 'classification'

    def load_and_preprocess_data(self, filepath):
        """
        Load and preprocess the student performance dataset
        """

        print("Loading dataset...")
        df = pd.read_csv(filepath)
        print(f"Dataset loaded with shape: {df.shape}")

        # Display basic info
        print("\nDataset Info:")
        print(f"Total students: {len(df)}")
        print(f"Features: {list(df.columns)}")

        # Check for missing values
        print(f"\nMissing values:\n{df.isnull().sum()}")

        # Handle missing values if any
        df = df.dropna() # Simple approach - drop rows with missing values

        # Analyze the Grade column to determine problem type
        print("\n" + "="*50)
        print("GRADE COLUMN ANALYSIS")
        print("="*50)

        grade_data = df['Grade']
        print(f"Grade data type: {grade_data.dtype}")
        print(f"Unique values: {grade_data.unique()[:20]}") # Show first 20 unique
        print(f"Number of unique values: {grade_data.nunique()}")

        # Check if grades are numeric or categorical
        try:
            # Try to convert to numeric
            numeric_grades = pd.to_numeric(grade_data, errors='coerce')
            if numeric_grades.notna().all():
                self.problem_type = 'regression'
                print("✓ Grades are numeric - using regression approach")
            else:
                self.problem_type = 'classification'
                print("✓ Grades are categorical - using classification approach")
        except:
            self.problem_type = 'classification'
            print("✓ Grades are categorical - using classification approach")

        return df

    def preprocess_features(self, df, is_training=True):
        """
        Preprocess features: encode categorical variables and scale numerical features
        """

        # Separate features and target
        X = df.drop(['Student_ID'], axis=1)
        y = df['Grade'] if 'Grade' in df.columns else None

```

```

# Identify categorical and numerical columns in features
categorical_cols = ['Sex', 'High_School_Type', 'Scholarship', 'Additional_W
                     'Sports_activity', 'Transportation']

# Find which categorical columns actually exist in the data
existing_categorical_cols = [col for col in categorical_cols if col in X.co]

# Find numerical columns (all remaining columns except Grade)
potential_numerical = [col for col in X.columns if col != 'Grade' and col n
numerical_cols = []

for col in potential_numerical:
    if pd.api.types.is_numeric_dtype(X[col]):
        numerical_cols.append(col)
    else:
        # If it's not numeric, add to categorical
        existing_categorical_cols.append(col)

print(f"\nCategorical features: {existing_categorical_cols}")
print(f"Numerical features: {numerical_cols}")

# Encode categorical variables
X_encoded = X.copy()

for col in existing_categorical_cols:
    if col in X.columns:
        if is_training:
            self.label_encoders[col] = LabelEncoder()
            # Handle any NaN values
            X[col] = X[col].fillna('Unknown')
            X_encoded[col] = self.label_encoders[col].fit_transform(X[col])
        else:
            # Handle unseen categories during prediction
            try:
                X[col] = X[col].fillna('Unknown')
                X_encoded[col] = self.label_encoders[col].transform(X[col])
            except ValueError:
                # If unseen category, set to -1
                X_encoded[col] = -1

# Handle target variable (Grade)
if y is not None:
    if self.problem_type == 'classification':
        # Encode categorical grades
        if is_training:
            self.grade_encoder = LabelEncoder()
            y_encoded = self.grade_encoder.fit_transform(y.astype(str))
            print(f"\nGrade categories: {list(self.grade_encoder.classes_)}")
        else:
            try:
                y_encoded = self.grade_encoder.transform(y.astype(str))
            except:
                y_encoded = -1 * np.ones(len(y))
    else:
        # Convert to float for regression
        y_encoded = pd.to_numeric(y, errors='coerce').fillna(0).astype(float)

```

```

    else:
        y_encoded = None

    # Scale numerical features
    if numerical_cols:
        if is_training:
            X_encoded[numerical_cols] = self.scaler.fit_transform(X_encoded[numerical_cols])
        else:
            # Ensure columns match and handle missing columns
            for col in numerical_cols:
                if col not in X_encoded.columns:
                    X_encoded[col] = 0
            X_encoded[numerical_cols] = self.scaler.transform(X_encoded[numerical_cols])

    # Store feature columns
    if is_training:
        self.feature_columns = X_encoded.columns.tolist()
    else:
        # Ensure all training columns exist
        for col in self.feature_columns:
            if col not in X_encoded.columns:
                X_encoded[col] = 0
        X_encoded = X_encoded[self.feature_columns]

    return X_encoded, y_encoded

def train_model(self, X, y, model_type='auto', test_size=0.2):
    """
    Train the appropriate model based on problem type
    """
    # Split the data
    X_train, X_test, y_train, y_test = train_test_split(
        X, y, test_size=test_size, random_state=42
    )

    # Auto-detect best model if 'auto' is selected
    if model_type == 'auto':
        if self.problem_type == 'regression':
            # Try multiple models and pick the best
            models_to_try = {
                'linear': LinearRegression(),
                'ridge': Ridge(alpha=1.0),
                'random_forest': RandomForestRegressor(n_estimators=100, random_state=42),
                'gradient_boosting': GradientBoostingRegressor(n_estimators=100)
            }
        else:
            models_to_try = {
                'logistic': LogisticRegression(max_iter=1000, random_state=42),
                'random_forest': RandomForestClassifier(n_estimators=100, random_state=42),
                'gradient_boosting': GradientBoostingClassifier(n_estimators=100)
            }

        best_score = -np.inf
        best_model_name = None
        best_model = None

```

```

print(f"\nTrying multiple {self.problem_type} models...")

for name, model in models_to_try.items():
    model.fit(X_train, y_train)
    if self.problem_type == 'regression':
        score = model.score(X_test, y_test)
    else:
        score = accuracy_score(y_test, model.predict(X_test))

    print(f"{name}: {score:.4f}")

    if score > best_score:
        best_score = score
        best_model_name = name
        best_model = model

self.model = best_model
print(f"\nSelected best model: {best_model_name} with score: {best_score:.4f}")

else:
    # Use specified model type
    if self.problem_type == 'regression':
        if model_type == 'linear':
            self.model = LinearRegression()
        elif model_type == 'ridge':
            self.model = Ridge(alpha=1.0)
        elif model_type == 'lasso':
            self.model = Lasso(alpha=0.01)
        elif model_type == 'random_forest':
            self.model = RandomForestRegressor(n_estimators=100, random_state=42)
        elif model_type == 'gradient_boosting':
            self.model = GradientBoostingRegressor(n_estimators=100, random_state=42)
        else:
            self.model = LinearRegression()
    else:
        if model_type == 'logistic':
            self.model = LogisticRegression(max_iter=1000, random_state=42)
        elif model_type == 'random_forest':
            self.model = RandomForestClassifier(n_estimators=100, random_state=42)
        elif model_type == 'gradient_boosting':
            self.model = GradientBoostingClassifier(n_estimators=100, random_state=42)
        else:
            self.model = LogisticRegression(max_iter=1000, random_state=42)

    print(f"\nTraining {model_type} {self.problem_type} model...")
    self.model.fit(X_train, y_train)

# Make predictions
y_pred_train = self.model.predict(X_train)
y_pred_test = self.model.predict(X_test)

# Evaluate model
self.evaluate_model(y_train, y_pred_train, y_test, y_pred_test)

self.is_trained = True
return self.model

```

```

def evaluate_model(self, y_train, y_pred_train, y_test, y_pred_test):
    """
    Evaluate model performance using appropriate metrics
    """
    print("\n" + "*50")
    print(f"MODEL EVALUATION METRICS ({self.problem_type.upper()})")
    print("*50")

    if self.problem_type == 'regression':
        # Regression metrics
        print("\nTRAINING SET:")
        print(f"MAE: {mean_absolute_error(y_train, y_pred_train):.4f}")
        print(f"RMSE: {np.sqrt(mean_squared_error(y_train, y_pred_train)):.4f}")
        print(f"R² Score: {r2_score(y_train, y_pred_train):.4f}")

        print("\nTESTING SET:")
        print(f"MAE: {mean_absolute_error(y_test, y_pred_test):.4f}")
        print(f"RMSE: {np.sqrt(mean_squared_error(y_test, y_pred_test)):.4f}")
        print(f"R² Score: {r2_score(y_test, y_pred_test):.4f}")

    else:
        # Classification metrics
        print("\nTRAINING SET:")
        print(f"Accuracy: {accuracy_score(y_train, y_pred_train):.4f}")

        print("\nTESTING SET:")
        print(f"Accuracy: {accuracy_score(y_test, y_pred_test):.4f}")

        # Get unique classes present in test set
        unique_classes = np.unique(y_test)

        if len(unique_classes) <= 10: # Only show detailed report if not too many
            print("\nClassification Report (Test Set):")
            print(classification_report(y_test, y_pred_test, zero_division=0))

            print("\nConfusion Matrix (Test Set):")
            print(confusion_matrix(y_test, y_pred_test))

def predict_grade(self, student_data):
    """
    Predict grade for a single student
    """
    if not self.is_trained:
        raise Exception("Model not trained yet. Please train the model first.")

    # Convert input to DataFrame
    if isinstance(student_data, dict):
        df_input = pd.DataFrame([student_data])
    else:
        df_input = student_data

    # Preprocess
    X_processed, _ = self.preprocess_features(df_input, is_training=False)

    # Predict

```

```

prediction = self.model.predict(X_processed)[0]

# Convert back to original grade format if classification
if self.problem_type == 'classification' and self.grade_encoder is not None
    # Handle float predictions (round to nearest integer for class index)
    if isinstance(prediction, (np.floating, float)):
        prediction = int(round(prediction))
    # Ensure prediction is within valid range
    if 0 <= prediction < len(self.grade_encoder.classes_):
        prediction = self.grade_encoder.inverse_transform([prediction])[0]
    else:
        prediction = "Unknown"

return prediction

def save_model(self, filepath='student_performance_model.pkl'):
    """
    Save the trained model and preprocessors
    """
    if not self.is_trained:
        raise Exception("No trained model to save.")

    model_data = {
        'model': self.model,
        'scaler': self.scaler,
        'label_encoders': self.label_encoders,
        'grade_encoder': self.grade_encoder,
        'feature_columns': self.feature_columns,
        'problem_type': self.problem_type
    }

    joblib.dump(model_data, filepath)
    print(f"\nModel saved successfully to {filepath}")

def load_model(self, filepath='student_performance_model.pkl'):
    """
    Load a trained model
    """
    if not os.path.exists(filepath):
        raise FileNotFoundError(f"Model file {filepath} not found.")

    model_data = joblib.load(filepath)
    self.model = model_data['model']
    self.scaler = model_data['scaler']
    self.label_encoders = model_data['label_encoders']
    self.grade_encoder = model_data.get('grade_encoder')
    self.feature_columns = model_data['feature_columns']
    self.problem_type = model_data.get('problem_type', 'classification')
    self.is_trained = True

    print(f"\nModel loaded successfully from {filepath}")
    print(f"Problem type: {self.problem_type}")

def analyze_feature_importance(self):
    """
    Analyze feature importance
    """

```

```

"""
if self.feature_columns is None:
    print("\nFeature information not available.")
    return

print("\n" + "*50)
print("FEATURE ANALYSIS")
print("*50)

if hasattr(self.model, 'feature_importances_'):
    # Tree-based models
    importances = self.model.feature_importances_
    indices = np.argsort(importances)[::-1]

    print("\nFeature Importance (higher = more important):")
    for i, idx in enumerate(indices):
        if i < min(10, len(self.feature_columns)):
            print(f"{i+1}. {self.feature_columns[idx]}: {importances[idx]}")

elif hasattr(self.model, 'coef_'):
    # Linear models
    coefficients = self.model.coef_
    if len(coefficients.shape) > 1:
        # For multi-class, take mean absolute coefficient
        coefficients = np.mean(np.abs(coefficients), axis=0)

    indices = np.argsort(np.abs(coefficients))[::-1]

    print("\nFeature Coefficients (absolute value):")
    for i, idx in enumerate(indices):
        if i < min(10, len(self.feature_columns)):
            print(f"{i+1}. {self.feature_columns[idx]}: {coefficients[idx]}")

else:
    print("Feature importance not available for this model type.")

def main():
"""
Main CLI interface for the Student Performance Prediction System
"""
predictor = StudentPerformancePredictor()

print("*60)
print("STUDENT PERFORMANCE PREDICTION SYSTEM")
print("*60)

while True:
    print("\n" + "-"*40)
    print("MAIN MENU")
    print("-"*40)
    print("1. Train new model")
    print("2. Load existing model")
    print("3. Predict student grade")
    print("4. Analyze feature importance")
    print("5. Exit")

    choice = input("\nEnter your choice (1-5): ").strip()

```

```

if choice == '1':
    # Train new model
    filepath = input("Enter dataset path (default: Students Performance .csv")
    if not filepath:
        filepath = "Students Performance .csv"

try:
    # Load and preprocess data
    df = predictor.load_and_preprocess_data(filepath)

    # Preprocess features
    X, y = predictor.preprocess_features(df, is_training=True)

    if predictor.problem_type == 'regression':
        print("\nSelect regression model:")
        print("1. Linear Regression")
        print("2. Ridge Regression")
        print("3. Lasso Regression")
        print("4. Random Forest")
        print("5. Gradient Boosting")
        print("6. Auto (try multiple and pick best)")
    else:
        print("\nSelect classification model:")
        print("1. Logistic Regression")
        print("2. Random Forest")
        print("3. Gradient Boosting")
        print("4. Auto (try multiple and pick best)")

    model_choice = input("Enter choice (1-6): ").strip()

    model_types_reg = {
        '1': 'linear',
        '2': 'ridge',
        '3': 'lasso',
        '4': 'random_forest',
        '5': 'gradient_boosting',
        '6': 'auto'
    }

    model_types_clf = {
        '1': 'logistic',
        '2': 'random_forest',
        '3': 'gradient_boosting',
        '4': 'auto'
    }

    if predictor.problem_type == 'regression':
        model_type = model_types_reg.get(model_choice, 'auto')
    else:
        model_type = model_types_clf.get(model_choice, 'auto')

    # Train model
    predictor.train_model(X, y, model_type=model_type)

    # Save model

```

```

        save_choice = input("\nSave model? (y/n): ").strip().lower()
        if save_choice == 'y':
            model_name = input("Enter model filename (default: student_mode
            if not model_name:
                model_name = "student_model.pkl"
            predictor.save_model(model_name)

        except Exception as e:
            print(f"\nError: {e}")
            import traceback
            traceback.print_exc()

    elif choice == '2':
        # Load existing model
        model_name = input("Enter model filename (default: student_model.pkl):
        if not model_name:
            model_name = "student_model.pkl"

        try:
            predictor.load_model(model_name)
        except Exception as e:
            print(f"\nError loading model: {e}")

    elif choice == '3':
        # Predict student grade
        if not predictor.is_trained:
            print("\nPlease train or load a model first!")
            continue

        print("\n" + "-"*40)
        print("ENTER STUDENT DETAILS FOR PREDICTION")
        print("-"*40)
        print("(Press Enter to skip optional fields)")

        try:
            # Collect student data
            student_data = {}

            # Required fields
            student_data['Student_Age'] = int(input("Student Age: "))

            # Categorical fields
            sex = input("Sex (Male/Female): ").strip()
            if sex:
                student_data['Sex'] = sex

            school_type = input("High School Type (Public/Private/Other): ")
            if school_type:
                student_data['High_School_Type'] = school_type

            scholarship = input("Scholarship (Yes/No): ").strip()
            if scholarship:
                student_data['Scholarship'] = scholarship

            additional_work = input("Additional Work (Yes/No): ").strip()
            if additional_work:

```

```

student_data['Additional_Work'] = additional_work

sports = input("Sports Activity (Yes/No): ").strip()
if sports:
    student_data['Sports_activity'] = sports

transport = input("Transportation (Bus/Walk/Other): ").strip()
if transport:
    student_data['Transportation'] = transport

# Numerical fields
study_hours = input("Weekly Study Hours: ").strip()
if study_hours:
    student_data['Weekly_Study_Hours'] = float(study_hours)

attendance = input("Attendance Percentage (0-100): ").strip()
if attendance:
    student_data['Attendance'] = float(attendance)

reading = input("Reading Score (0-100): ").strip()
if reading:
    student_data['Reading'] = float(reading)

notes = input("Notes Taking Score (0-100): ").strip()
if notes:
    student_data['Notes'] = float(notes)

listening = input("Listening in Class Score (0-100): ").strip()
if listening:
    student_data['Listening_in_Class'] = float(listening)

project = input("Project Work Score (0-100): ").strip()
if project:
    student_data['Project_work'] = float(project)

# Make prediction
predicted_grade = predictor.predict_grade(student_data)

print("\n" + "="*40)
print("PREDICTION RESULT")
print("="*40)
print(f"Predicted Grade: {predicted_grade}")

# Provide interpretation for numeric grades
if predictor.problem_type == 'regression':
    try:
        grade_num = float(predicted_grade)
        if grade_num >= 90:
            print("Performance: Excellent (A)")
        elif grade_num >= 80:
            print("Performance: Very Good (B)")
        elif grade_num >= 70:
            print("Performance: Good (C)")
        elif grade_num >= 60:
            print("Performance: Satisfactory (D)")
    else:

```

```

        print("Performance: Needs Improvement (F)")
    except:
        pass

    except Exception as e:
        print(f"\nError in prediction: {e}")
        import traceback
        traceback.print_exc()

    elif choice == '4':
        # Analyze feature importance
        if not predictor.is_trained:
            print("\nPlease train or load a model first!")
            continue

        predictor.analyze_feature_importance()

    elif choice == '5':
        print("\nThank you for using Student Performance Prediction System!")
        break

    else:
        print("\nInvalid choice. Please try again.")

if __name__ == "__main__":
    main()

```

=====
STUDENT PERFORMANCE PREDICTION SYSTEM
=====

MAIN MENU

1. Train new model
2. Load existing model
3. Predict student grade
4. Analyze feature importance
5. Exit

Invalid choice. Please try again.

MAIN MENU

1. Train new model
2. Load existing model
3. Predict student grade
4. Analyze feature importance
5. Exit

Invalid choice. Please try again.

MAIN MENU

1. Train new model
2. Load existing model
3. Predict student grade
4. Analyze feature importance
5. Exit

Invalid choice. Please try again.

MAIN MENU

1. Train new model
2. Load existing model
3. Predict student grade
4. Analyze feature importance
5. Exit

Invalid choice. Please try again.

MAIN MENU

1. Train new model
2. Load existing model
3. Predict student grade
4. Analyze feature importance
5. Exit

Invalid choice. Please try again.

MAIN MENU

1. Train new model
2. Load existing model
3. Predict student grade
4. Analyze feature importance
5. Exit

Invalid choice. Please try again.

MAIN MENU

1. Train new model
2. Load existing model
3. Predict student grade
4. Analyze feature importance
5. Exit

Loading dataset...

Dataset loaded with shape: (145, 15)

Dataset Info:

Total students: 145

Features: ['Student_ID', 'Student_Age', 'Sex', 'High_School_Type', 'Scholarship', 'Additional_Work', 'Sports_activity', 'Transportation', 'Weekly_Study_Hours', 'Attendance', 'Reading', 'Notes', 'Listening_in_Class', 'Project_work', 'Grade']

Missing values:

Student_ID	0
Student_Age	0
Sex	0
High_School_Type	0
Scholarship	1
Additional_Work	0
Sports_activity	0
Transportation	0
Weekly_Study_Hours	0
Attendance	0
Reading	0
Notes	0
Listening_in_Class	0
Project_work	0
Grade	0

dtype: int64

=====
GRADE COLUMN ANALYSIS
=====

Grade data type: object

Unique values: ['AA' 'BA' 'CC' 'Fail' 'BB' 'CB' 'DD' 'DC']

Number of unique values: 8

✓ Grades are categorical - using classification approach

Categorical features: ['Sex', 'High_School_Type', 'Scholarship', 'Additional_Work', 'Sports_activity', 'Transportation', 'Student_Age', 'Attendance', 'Reading', 'Notes', 'Listening_in_Class', 'Project_work']

Numerical features: ['Weekly_Study_Hours']

Grade categories: ['AA', 'BA', 'BB', 'CB', 'CC', 'DC', 'DD', 'Fail']

Select classification model:

1. Logistic Regression
2. Random Forest
3. Gradient Boosting
4. Auto (try multiple and pick best)

```

Traceback (most recent call last):
  File "C:\Users\ASIM ALI\AppData\Local\Temp\ipykernel_29208\2537903796.py", line 46
7, in main
    predictor.train_model(X, y, model_type=model_type)
  File "C:\Users\ASIM ALI\AppData\Local\Temp\ipykernel_29208\2537903796.py", line 23
5, in train_model
    self.model.fit(X_train, y_train)
  File "C:\Users\ASIM ALI\anaconda3\envs\fresh_env\lib\site-packages\sklearn\base.p
y", line 1365, in wrapper
    return fit_method(estimator, *args, **kwargs)
  File "C:\Users\ASIM ALI\anaconda3\envs\fresh_env\lib\site-packages\sklearn\ensembl
e_forest.py", line 359, in fit
    X, y = validate_data(
  File "C:\Users\ASIM ALI\anaconda3\envs\fresh_env\lib\site-packages\sklearn\utils\v
alidation.py", line 2971, in validate_data
    X, y = check_X_y(X, y, **check_params)
  File "C:\Users\ASIM ALI\anaconda3\envs\fresh_env\lib\site-packages\sklearn\utils\v
alidation.py", line 1368, in check_X_y
    X = check_array(
  File "C:\Users\ASIM ALI\anaconda3\envs\fresh_env\lib\site-packages\sklearn\utils\v
alidation.py", line 1053, in check_array
    array = _asarray_with_order(array, order=order, dtype=dtype, xp=xp)
  File "C:\Users\ASIM ALI\anaconda3\envs\fresh_env\lib\site-packages\sklearn\utils\_a
rray_api.py", line 757, in _asarray_with_order
    array = numpy.asarray(array, order=order, dtype=dtype)
  File "C:\Users\ASIM ALI\anaconda3\envs\fresh_env\lib\site-packages\pandas\core\gen
eric.py", line 2168, in __array__
    arr = np.asarray(values, dtype=dtype)
ValueError: could not convert string to float: 'BA'
Training random_forest classification model...

```

Error: could not convert string to float: 'BA'

MAIN MENU

1. Train new model
2. Load existing model
3. Predict student grade
4. Analyze feature importance
5. Exit

Please train or load a model first!

MAIN MENU

1. Train new model
2. Load existing model
3. Predict student grade
4. Analyze feature importance
5. Exit

Please train or load a model first!

MAIN MENU

1. Train new model
2. Load existing model
3. Predict student grade
4. Analyze feature importance
5. Exit

Invalid choice. Please try again.

MAIN MENU

1. Train new model
2. Load existing model
3. Predict student grade
4. Analyze feature importance
5. Exit

Invalid choice. Please try again.

MAIN MENU

1. Train new model
2. Load existing model
3. Predict student grade
4. Analyze feature importance
5. Exit

Thank you for using Student Performance Prediction System!

In [9]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from collections import Counter

def explore_dataset(filepath="Students Performance .csv"):
    """
    Explore and visualize the student performance dataset
    """
    # Load data
    df = pd.read_csv(filepath)

    print("*"*60)
    print("STUDENT PERFORMANCE DATASET EXPLORATION")
    print("*"*60)

    # Basic statistics
    print("\n1. DATASET OVERVIEW")
    print("-"*40)
    print(f"Shape: {df.shape}")
    print(f"Columns: {list(df.columns)}")
    print(f"\nData Types:\n{df.dtypes}")
```

```

# Check for missing values
print("\n2. MISSING VALUES")
print("-"*40)
print(df.isnull().sum())

# Grade analysis
print("\n3. GRADE DISTRIBUTION")
print("-"*40)

grade_data = df['Grade']

# Try to determine if grades are numeric or categorical
try:
    numeric_grades = pd.to_numeric(grade_data, errors='coerce')
    if numeric_grades.notna().all():
        print("Grades are NUMERIC")
        print(f"Mean Grade: {grade_data.mean():.2f}")
        print(f"Median Grade: {grade_data.median():.2f}")
        print(f"Std Deviation: {grade_data.std():.2f}")
        print(f"Min Grade: {grade_data.min()}")
        print(f"Max Grade: {grade_data.max()}")
    else:
        print("Grades are CATEGORICAL")
        grade_counts = grade_data.value_counts()
        print(f"Unique grade values: {grade_data.nunique()}")
        print("\nGrade frequency:")
        for grade, count in grade_counts.head(10).items():
            print(f" {grade}: {count} ({count/len(df)*100:.1f}%)")
except:
    print("Grades are CATEGORICAL")
    grade_counts = grade_data.value_counts()
    print(f"Unique grade values: {grade_data.nunique()}")
    print("\nGrade frequency (top 10):")
    for grade, count in grade_counts.head(10).items():
        print(f" {grade}: {count} ({count/len(df)*100:.1f}%)")

# Summary statistics for numerical columns
print("\n4. NUMERICAL FEATURES SUMMARY")
print("-"*40)
numerical_cols = df.select_dtypes(include=[np.number]).columns
if len(numerical_cols) > 0:
    print(df[numerical_cols].describe())
else:
    print("No numerical columns found (excluding Grade)")

# Categorical columns analysis
print("\n5. CATEGORICAL FEATURES SUMMARY")
print("-"*40)
categorical_cols = df.select_dtypes(include=['object']).columns
categorical_cols = [col for col in categorical_cols if col != 'Grade']

for col in categorical_cols:
    print(f"\n{col}:")
    value_counts = df[col].value_counts()
    for val, count in value_counts.head(5).items():
        print(f" {val}: {count} ({count/len(df)*100:.1f}%)")

```

```

# Create visualizations
print("\n6. GENERATING VISUALIZATIONS...")

# Determine grid size based on number of columns
n_cols = min(3, len(df.columns))
n_rows = (len(df.columns) + n_cols - 1) // n_cols

fig, axes = plt.subplots(n_rows, n_cols, figsize=(5*n_cols, 4*n_rows))
if n_rows == 1:
    axes = [axes]
axes_flat = [ax for row in axes for ax in row]

plot_idx = 0

for col in df.columns:
    if col == 'Student_ID':
        continue

    ax = axes_flat[plot_idx]

    if col == 'Grade':
        # Special handling for Grade column
        try:
            numeric_grades = pd.to_numeric(df[col], errors='coerce')
            if numeric_grades.notna().all():
                ax.hist(df[col].dropna(), bins=20, edgecolor='black', alpha=0.7)
                ax.set_title(f'{col} Distribution')
                ax.set_xlabel(col)
            else:
                # Categorical grade - show bar chart
                grade_counts = df[col].value_counts().head(10)
                ax.bar(range(len(grade_counts)), grade_counts.values)
                ax.set_xticks(range(len(grade_counts)))
                ax.set_xticklabels(grade_counts.index, rotation=45, ha='right')
                ax.set_title(f'{col} Distribution (Top 10)')
        except:
            grade_counts = df[col].value_counts().head(10)
            ax.bar(range(len(grade_counts)), grade_counts.values)
            ax.set_xticks(range(len(grade_counts)))
            ax.set_xticklabels(grade_counts.index, rotation=45, ha='right')
            ax.set_title(f'{col} Distribution (Top 10)')

    elif pd.api.types.is_numeric_dtype(df[col]):
        # Numerical column - histogram
        ax.hist(df[col].dropna(), bins=20, edgecolor='black', alpha=0.7)
        ax.set_title(f'{col} Distribution')
        ax.set_xlabel(col)

    else:
        # Categorical column - bar chart
        value_counts = df[col].value_counts().head(8)
        ax.bar(range(len(value_counts)), value_counts.values)
        ax.set_xticks(range(len(value_counts)))
        ax.set_xticklabels(value_counts.index, rotation=45, ha='right')
        ax.set_title(f'{col} Distribution')

```

```
        ax.set_ylabel('Frequency')
        plot_idx += 1

    # Hide unused subplots
    for idx in range(plot_idx, len(axes_flat)):
        axes_flat[idx].set_visible(False)

    plt.tight_layout()
    plt.savefig('student_performance_analysis.png', dpi=300, bbox_inches='tight')
    plt.show()

    print("\nVisualizations saved to 'student_performance_analysis.png'")


# Correlation analysis for numerical features
print("\n7. CORRELATION ANALYSIS")
print("-"*40)
numerical_cols = df.select_dtypes(include=[np.number]).columns.tolist()
if 'Student_ID' in numerical_cols:
    numerical_cols.remove('Student_ID')

if len(numerical_cols) > 1:
    corr_matrix = df[numerical_cols].corr()
    print("Correlation matrix:")
    print(corr_matrix)

    # Plot correlation heatmap
    plt.figure(figsize=(10, 8))
    sns.heatmap(corr_matrix, annot=True, fmt='.2f', cmap='coolwarm', center=0)
    plt.title('Feature Correlation Heatmap')
    plt.tight_layout()
    plt.savefig('correlation_heatmap.png', dpi=300, bbox_inches='tight')
    plt.show()
    print("\nCorrelation heatmap saved to 'correlation_heatmap.png'")


return df

if __name__ == "__main__":
    df = explore_dataset()
```

STUDENT PERFORMANCE DATASET EXPLORATION

1. DATASET OVERVIEW

Shape: (145, 15)

Columns: ['Student_ID', 'Student_Age', 'Sex', 'High_School_Type', 'Scholarship', 'Additional_Work', 'Sports_activity', 'Transportation', 'Weekly_Study_Hours', 'Attendance', 'Reading', 'Notes', 'Listening_in_Class', 'Project_work', 'Grade']

Data Types:

Student_ID	object
Student_Age	object
Sex	object
High_School_Type	object
Scholarship	object
Additional_Work	object
Sports_activity	object
Transportation	object
Weekly_Study_Hours	int64
Attendance	object
Reading	object
Notes	object
Listening_in_Class	object
Project_work	object
Grade	object
dtype:	object

2. MISSING VALUES

Student_ID	0
Student_Age	0
Sex	0
High_School_Type	0
Scholarship	1
Additional_Work	0
Sports_activity	0
Transportation	0
Weekly_Study_Hours	0
Attendance	0
Reading	0
Notes	0
Listening_in_Class	0
Project_work	0
Grade	0
dtype:	int64

3. GRADE DISTRIBUTION

Grades are CATEGORICAL

Unique grade values: 8

Grade frequency:

AA: 35 (24.1%)

BA: 24 (16.6%)

BB: 21 (14.5%)
CC: 17 (11.7%)
DD: 17 (11.7%)
DC: 13 (9.0%)
CB: 10 (6.9%)
Fail: 8 (5.5%)

4. NUMERICAL FEATURES SUMMARY

Weekly_Study_Hours

count	145.000000
mean	2.331034
std	4.249273
min	0.000000
25%	0.000000
50%	0.000000
75%	2.000000
max	12.000000

5. CATEGORICAL FEATURES SUMMARY

Student_ID:

STUDENT1: 1 (0.7%)
STUDENT74: 1 (0.7%)
STUDENT94: 1 (0.7%)
STUDENT95: 1 (0.7%)
STUDENT96: 1 (0.7%)

Student_Age:

19-22: 70 (48.3%)
18: 65 (44.8%)
23-27: 10 (6.9%)

Sex:

Male: 87 (60.0%)
Female: 58 (40.0%)

High_School_Type:

State: 103 (71.0%)
Private: 25 (17.2%)
Other: 17 (11.7%)

Scholarship:

50%: 76 (52.4%)
75%: 42 (29.0%)
100%: 23 (15.9%)
25%: 3 (2.1%)

Additional_Work:

No: 96 (66.2%)
Yes: 49 (33.8%)

Sports_activity:

No: 87 (60.0%)
Yes: 58 (40.0%)

Transportation:

Private: 84 (57.9%)
Bus: 61 (42.1%)

Attendance:

Always: 98 (67.6%)
Sometimes: 25 (17.2%)
Never: 21 (14.5%)
3: 1 (0.7%)

Reading:

No: 76 (52.4%)
Yes: 69 (47.6%)

Notes:

Yes: 77 (53.1%)
No: 66 (45.5%)
6: 2 (1.4%)

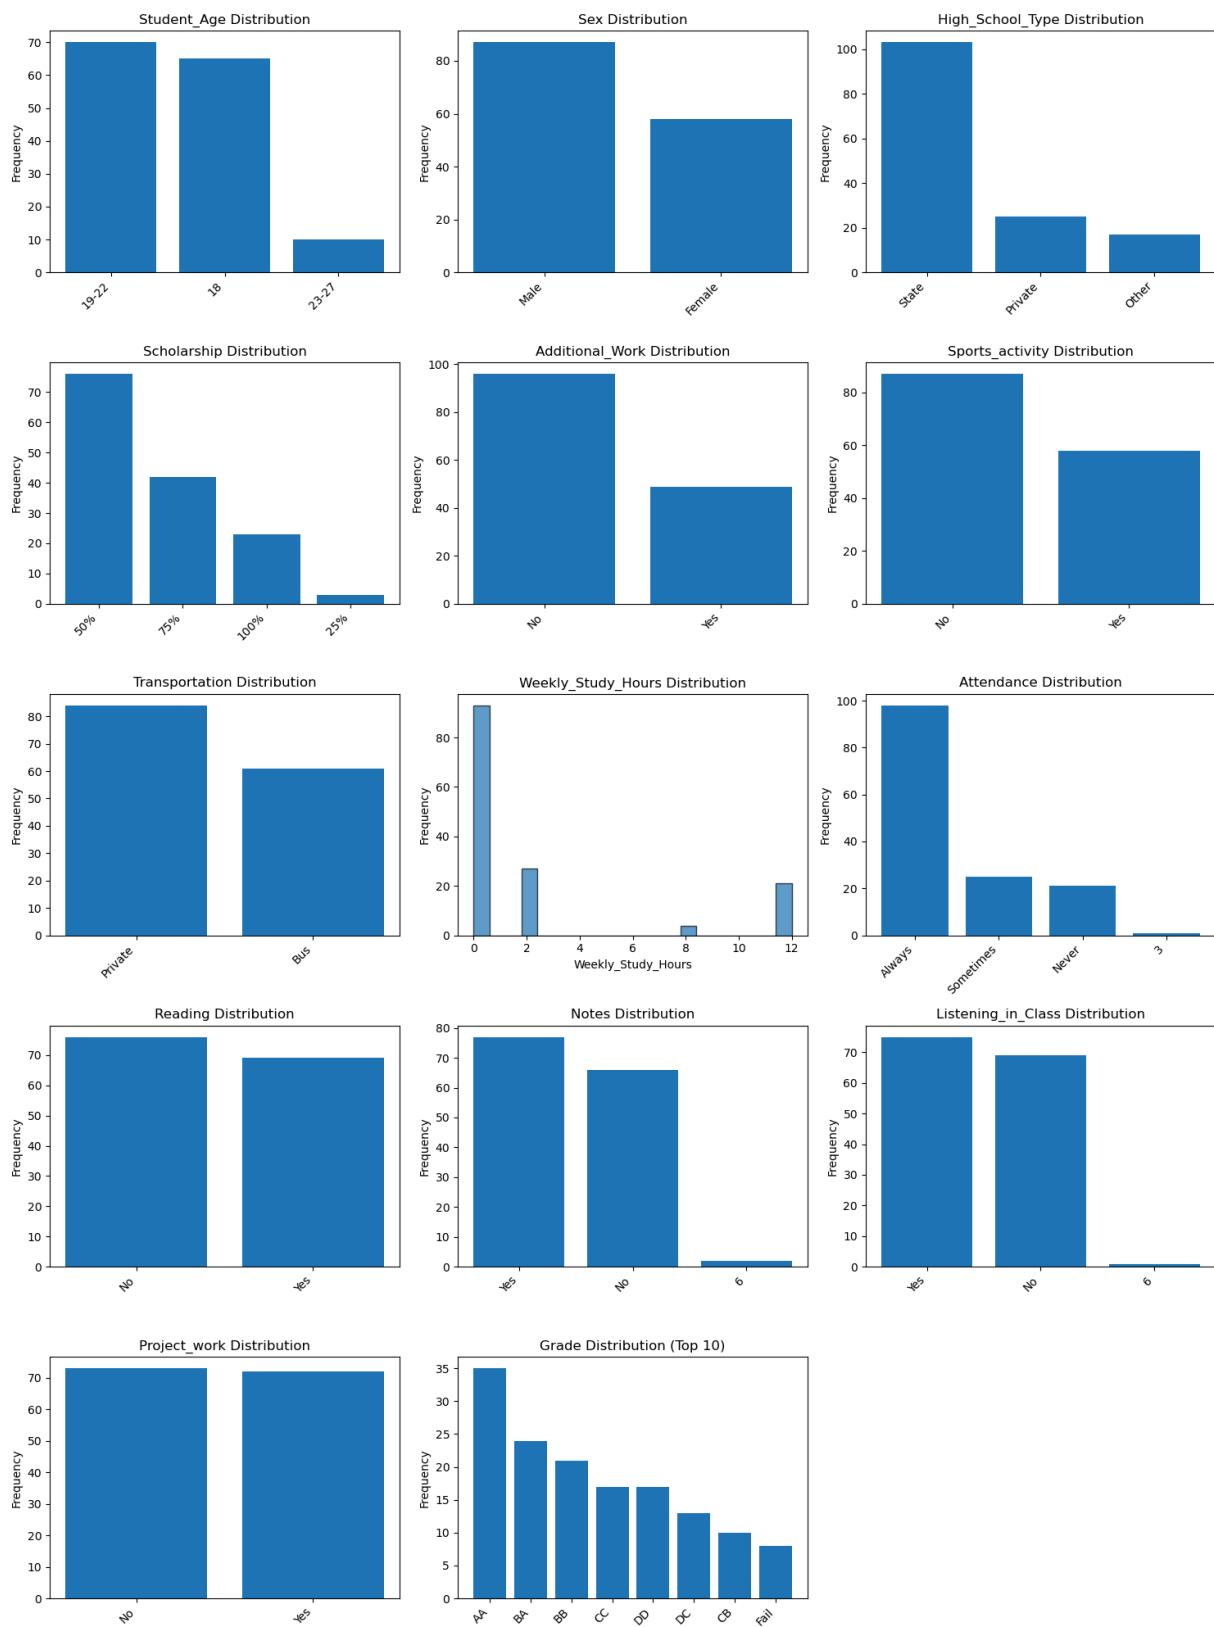
Listening_in_Class:

Yes: 75 (51.7%)
No: 69 (47.6%)
6: 1 (0.7%)

Project_work:

No: 73 (50.3%)
Yes: 72 (49.7%)

6. GENERATING VISUALIZATIONS...



Visualizations saved to 'student_performance_analysis.png'

7. CORRELATION ANALYSIS
